

MIRAGE · Demo Script

A privacy firewall for AI: personal details become placeholders and secrets never leave the device, before the prompt is sent.

Team Nexora **Date** 30 September 2026 **Length** about 10 minutes + Q&A

Parts Browser extension · Desktop companion · CLI

All data in this script is made up. Never type real details on stage. Every input below was run through MIRAGE's detector on 29 Sep 2026, and the scores and outputs shown are what it produced.

Running order

TIME	PART	WHO	WHAT THE AUDIENCE SEES
0:00	Opening	Presenter	The problem: people paste PAN, Aadhaar and API keys into AI chats.
0:45	A · Browser extension	Darshan	ChatGPT website: clean prompt passes, personal data gets hidden, secrets get blocked, Private Compose.
4:00	B · Desktop companion	Darshan	ChatGPT Mac app: "MIRAGE active" badge, the review opens as you type, Protect & Send, secret removed.
7:00	C · CLI	Ninad	Terminal: stop a secret in a log, hide details, put them back in the AI's reply, chat with Claude.
9:15	Closing	Presenter	One engine, three products, all on the device. Then Q&A.

Before you present (15 minutes earlier)

Browser (Chrome)

- ❑ MIRAGE extension loaded and pinned (shield in the toolbar shows *Protection on*).
- ❑ chatgpt.com logged in, **new chat** open.
- ❑ Turn off other extensions (Grammarly etc.).
- ❑ Zoom the page to 125% so the room can read it.

Desktop (Mac)

- ❑ MIRAGE running: shield in the menu bar, status *Protected* for ChatGPT.
- ❑ ChatGPT desktop app open on a **new chat**; the **"MIRAGE active"** badge shows above the message box.
- ❑ Menu bar → MIRAGE: *Show "MIRAGE active"* is on.
- ❑ If the badge doesn't show: System Settings → Privacy & Security → Accessibility → MIRAGE on,

CLI (terminal)

- ❑ Installed: `npm install -g https://github.com/NinadPandith/Mirage-Cli/raw/main/release/nexora-mirage-cli-0.1.8.tgz`
- ❑ Reset so the demo map is the latest: `mirage maps clear --yes`
- ❑ Practice files: `mirage demo` then `cd mirage-demo`
- ❑ Terminal font 20 pt or more.
- ❑ For the chat step: `claude -p "say hi"` answers (Claude Code logged in).

General

- ❑ Phone hotspot ready in case the Wi-Fi is bad.
- ❑ This PDF open on a second screen, or the inputs (page 6) in a notes app to copy from.
- ❑ Notifications off (Do Not Disturb).

then relaunch MIRAGE.

Opening · 0:00–0:45

"Every day people paste their PAN, Aadhaar, phone numbers and even API keys into ChatGPT and Claude. Once it's sent, it's on someone else's server. MIRAGE is a privacy firewall for AI: it checks every prompt **on your own device**, replaces personal details with placeholders, removes secrets completely, and nothing is sent until you choose. The same engine runs in three places: the browser, the ChatGPT and Claude desktop apps, and the terminal."

Part A · Browser extension on chatgpt.com · 0:45–4:00

A1 · A normal prompt passes untouched (20 s)

TYPE IN CHATGPT

Explain how UPI payments work in three simple lines.

WHAT HAPPENS

No popup. The message sends as usual. **Safe · 0/100**

SAY

"No false alarms: if there's nothing sensitive, MIRAGE stays out of the way."

A2 · Personal details get hidden (60 s) ★ main moment

TYPE (OR PASTE) IN CHATGPT, THEN PRESS ENTER

Hi, I'm Priya Nair. My PAN is BNZPM2501K, Aadhaar 2341 2341 2346, phone 98450 12345 and email priya.nair@gmail.com. Write a short letter to my bank asking them to update my KYC.

WHAT HAPPENS

- While typing, the shield badge by the message box turns red.
- Enter is stopped and the review opens: **5 items will be hidden before sending** · **Critical · 100/100** · Name, PAN, Aadhaar, Phone, Email.
- Click **Review** to show "Why 100?", the points behind the score.
- Click **Protect & send**. ChatGPT receives:

Hi, I'm «PERSON_1». My PAN is «PAN_1», Aadhaar «AADHAAR_1», phone «PHONE_1» and email «EMAIL_1». Write a short letter to my bank asking them to update my KYC.

- In ChatGPT's reply, **hover a placeholder**: MIRAGE shows the real value, on this laptop only. Use **Copy with details** under the reply to copy the letter with the real details filled in.

SAY

"ChatGPT never saw Priya's PAN or Aadhaar, only placeholders. The real values stay encrypted on this laptop and come back in the reply only here. The Aadhaar is confirmed with its checksum, so random 12-digit numbers don't trigger it."

A3 · Secrets are blocked, not hidden (40 s)

TYPE, THEN PRESS ENTER

```
My deploy keeps failing. Config: AWS_ACCESS_KEY_ID=AKIAQ7Z3MIRAGEDEM042 and db
password=S3cr3tPass@2026. What is wrong?
```

WHAT HAPPENS

Critical · 100/100 · AWS access key + password. Protect & send sends:

```
My deploy keeps failing. Config: AWS_ACCESS_KEY_ID=«AWS_ACCESS_KEY_REMOVED» and db
password=«PASSWORD_REMOVED». What is wrong?
```

Point out: **Send anyway** asks for a second confirmation for secrets.

SAY

"Keys and passwords are never stored and never put back. They're simply removed. One leaked AWS key can cost lakhs."

A4 · Health data: keep what the AI needs (30 s)

```
My name is Arjun Menon, DOB 14/03/2002. My HbA1c is 7.9 and I take metformin. What diet
should I follow?
```

High · 52/100 · the name and DOB are hidden; **HbA1c 7.9** and **metformin** are kept (the AI needs them) but counted in the score, with +15 for health data tied to a person.

```
My name is «PERSON_1», DOB «DOB_1». My HbA1c is 7.9 and I take metformin. What diet should I
follow?
```

"It's not a blind filter: the medical detail stays so the answer is still useful, but the AI can't tell whose it is."

A5 · Tricks don't work (20 s, optional)

```
Why was my payment declined? Card 4111 1111 1111 1111, expiry 12/28, OTP 482913.
```

Critical · 99/100 → Card «**CARD_REMOVED**», OTP «**OTP_REMOVED**» (the card is confirmed with the Luhn check).

```
Email me at priya.nair [at] gmail [dot] com and call +91-98450-12345.
```

The disguised email and the dashed phone number are both caught: «**EMAIL_1**», «**PHONE_1**».

A6 · Private Compose: the strongest mode (40 s)

1. Click the MIRAGE shield in the toolbar → **Private Compose**. A side panel opens.
2. Paste the A2 prompt into the side panel. It shows **5 items will be hidden or removed** and a live preview with the placeholders.
3. Click **Insert protected**. Only the protected text lands in ChatGPT's box → press Send.

"Here ChatGPT's page never sees the raw text at all, not even while you type."

Part B · Desktop companion in the ChatGPT Mac app · 4:00–7:00

"Many people don't use the website at all; they use the ChatGPT and Claude desktop apps. A browser extension can't see those, so we built a desktop companion that works the same way."

B1 · It's on (15 s)

Point at the **shield in the menu bar** (status: ChatGPT *Protected*) and the green "**MIRAGE active**" badge above ChatGPT's message box.

"This badge tells you MIRAGE is watching this box. It only reads the message box of ChatGPT and Claude: no keylogger, no screen recording."

B2 · The review opens as you type (60 s) ★ main moment

TYPE IN THE CHATGPT APP (DON'T PRESS ENTER YET)

I'm Rahul Sharma. My PAN is PQRSX6789K and my UPI is rahul.sharma@okhdfcbank. Draft an email to HDFC support, IFSC HDFC0001234, asking why my refund is late.

WHAT HAPPENS

- The badge turns amber: **4 items will be hidden**.
- About 0.15 s after you stop typing, the review opens above the box: **High privacy risk** · 55/100 · "Your information has NOT been sent yet." · Name, PAN, UPI, IFSC.
- Press **Enter** to show that nothing is sent while the review is open.
- Click **Protect & Send**. MIRAGE rewrites the box, sends it and shows "Sent protected · 4 hidden · 0 removed". The chat shows:

I'm «PERSON_1». My PAN is «PAN_1» and my UPI is «UPI_1». Draft an email to HDFC support, IFSC «IFSC_1», asking why my refund is late.

"Exactly like the extension, but inside a native app. The message is held until you decide."

B3 · A secret in code (40 s)

```
Fix this code: const key = "sk-proj-Xy7Qm2Lp9Rt4Vb8Nc3Kd6Hf1Js5Gw0Ez"; fetch(url, { headers: { Authorization: key } })
```

The badge shows **Secret · won't be sent**; the review says **Critical privacy risk** · 71/100 . Protect & Send sends:

```
Fix this code: const key = "«OPENAI_KEY_REMOVED»"; fetch(url, { headers: { Authorization: key } })
```

B4 · Cancel leaves everything untouched (20 s)

Type the A2 prompt, and when the review opens click **Cancel**. The text stays in the box exactly as typed, and nothing was sent.

B5 · Private Compose on the desktop (30 s, optional)

Press `⌘M` (or menu bar → **Private Compose**). Type the B2 prompt in MIRAGE's own window, check the preview, choose ChatGPT, then click **Insert protected**. Only the protected text goes into ChatGPT.

Windows: the same companion is available for Windows as a beta. Install it with **Install-MIRAGE-Desktop.bat** from github.com/Darshanj-dev/Mirage#download. It is built and self-tested on Windows by GitHub Actions but **not yet tried with the ChatGPT or Claude apps on Windows**, so say "beta" if asked.

Part C · CLI for developers (Ninad) · 7:00–9:15

"Same engine, now for developers who paste logs and code into AI." All outputs below are real, captured on 29 Sep 2026 with *mirage-cli* 0.1.8.

C1 · Show the log (10 s)

```
cat error.log
```

```
2026-09-29 10:02:12 DEBUG loaded env OPENAI_API_KEY=sk-proj-Xy7Qm2Lp9Rt4Vb8Nc3Kd6Hf1Js5Gw0Ez
2026-09-29 10:02:14 ERROR refund failed for customer priya.nair@example-mail.in
2026-09-29 10:02:14 ERROR kyc lookup: aadhaar 2341 2341 2346 not verified
2026-09-29 10:02:14 ERROR callback to +91 98450 12345 timed out after 30000ms
```

"A normal server log: an API key, an email, an Aadhaar number and a phone number. Developers paste this into ChatGPT every day."

C2 · It stops secrets (15 s)

```
mirage mask error.log
```

```
Line 2 looks like an API key (sk-.....w0Ez). Nothing was written.
Remove it, or run again with --strip-secrets.
```

C3 · Remove the secret, hide the rest (20 s)

```
mirage mask error.log --strip-secrets
```

```
... OPENAI_API_KEY=«SECRET_REMOVED»
... refund failed for customer «EMAIL_1»
... kyc lookup: aadhaar «AADHAAR_1» not verified
... callback to «PHONE_1» timed out after 30000ms
1 secret removed. It is not saved anywhere.
3 items hidden · map m_20260929_... · run "mirage unmask" on the AI's reply
```

"This is what you paste into any AI."

C4 · The real details come back (20 s)

```
mirage unmask reply.txt
```

```
The refund for priya.nair@example-mail.in failed because the KYC check rejected Aadhaar 2341
2341 2346
before the payment step ran, so the callback to +91 98450 12345 never got a response ...
3 placeholders restored from map m_20260929_...
```

C5 · Scan a whole folder (15 s, optional)

```
mirage scan .
```

```
error.log
 2:53 API_KEY  sk-.....w0Ez      secret
 3:54 EMAIL    p.....@example-mail.in  hide
 4:49 AADHAAR   ....  .... 2346         hide
 5:41 PHONE    .....2345             hide
3 files checked · 3 to hide · 1 secret
```

C6 · Chat with Claude through MIRAGE (40 s, needs internet)

```
mirage chat claude
```

AT THE YOU > PROMPT TYPE

```
My PAN is PQRSX6789K. Write one line confirming my PAN.
```

```
sent > My PAN is «PAN_1». Write one line confirming my PAN.    ← what Claude received
claude >
Your PAN is PQRSX6789K.                                         ← real value put back
1 placeholder restored
```

Then type my password is Hunter@123 → *Not sent: it looks like a password (.....)*. Then /exit.

"Claude only ever saw the placeholder, and we can prove it from Claude Code's own saved history:"

```
grep -r "PQRSX6789K" ~/.claude/projects/*mirage-demo* | wc -l      # expected 0
grep -r "PAN_1" ~/.claude/projects/*mirage-demo* | wc -l          # expected 1 or more
```

```
Windows (PowerShell): (Select-String -Path "$HOME\.claude\projects\*mirage-demo*\*.jsonl" -
Pattern "PQRSX6789K" -SimpleMatch).Count
```

Closing · 9:15–10:00

"One detection engine and three products: the browser, the desktop apps and the terminal. Everything runs on your device. There's no MIRAGE server, no account, and secrets are never stored. It catches Aadhaar and cards with real checksums, scores the risk from 0 to 100, and explains every point. You stay in control: nothing is sent until you choose."

Q&A cheat sheet

QUESTION	ANSWER
Does my data go to your server?	No. Detection runs on the device (browser, Mac app, terminal). There is no MIRAGE server. Only the protected text goes to the AI.
How does it detect?	Layers: text clean-up (hidden characters, look-alike letters, Indian-script digits) → patterns with checksums (Aadhaar Verhoeff, card Luhn, PAN format) → entropy for random-looking keys → context words ("PAN", "DOB", "my name is") → risk score.
How accurate is it?	On our 181-prompt test set: precision 1.000, recall 0.996, F1 0.998, 0 false alarms on 28 clean look-alikes. Be honest: we wrote the test set ourselves; a blind test by others is next.
How fast?	0.05 ms per prompt on average for detection. On the Mac the review opens about 0.15 s after you stop typing.
What does it miss?	Names only with context ("my name is ..."); a PAN split with spaces (e.g. "BNZPM 2501 K"); non-English text. Details it misses are sent as typed. That's why the review shows exactly what will be hidden.
Is the desktop app a keylogger?	No. It reads only the message box of ChatGPT and Claude, only while they're in front. It watches only Return and clicks on Send, never records the screen and never logs text (its log has reason codes and timings only).
Where are the real values kept?	Encrypted on the device (AES-GCM) for 24 hours so replies can show them; secrets are never stored at all.
What about the reply?	The extension checks the AI's reply and shows the real values on hover, on your device only. The CLI's unmask does the same for text.
Which platforms?	Extension: ChatGPT, Gemini, Claude, Copilot, Perplexity (Chrome, Edge, Firefox). Desktop: ChatGPT and Claude apps on Mac; Windows is a beta. CLI: Windows, macOS, Linux.

If something goes wrong on stage

PROBLEM	DO THIS
No review on chatgpt.com	Check the toolbar shield says <i>Protection on</i> ; reload the ChatGPT tab (Cmd/Ctrl+R) and try again.
No badge or review in the ChatGPT app	Menu bar → MIRAGE → check ChatGPT shows <i>Protected</i> . Quit and reopen MIRAGE. Fallback: use Private Compose (⌘⌘M).
Wi-Fi is down	Switch to the hotspot. The CLI steps C1–C5 work offline; skip C6.
ChatGPT is slow to answer	Don't wait for the reply: the point is what was sent , which is visible at once in the chat.
mirage chat claude errors	Skip C6 and show C4 (unmask) instead; it tells the same story offline.

All inputs in one place (copy from here)

STEP	INPUT	EXPECTED
A1	Explain how UPI payments work in three simple lines.	0 · safe no popup
A2 / B4	Hi, I'm Priya Nair. My PAN is BNZPM2501K, Aadhaar 2341 2341 2346, phone 98450 12345 and email priya.nair@gmail.com. Write a short letter to my bank asking them to update my KYC.	100 · critical 5 hidden
A3	My deploy keeps failing. Config: AWS_ACCESS_KEY_ID=AKIAQ7Z3MIRAGEDEM042 and db password=S3cr3tPass@2026. What is wrong?	100 · critical 2 removed
A4	My name is Arjun Menon, DOB 14/03/2002. My HbA1c is 7.9 and I take metformin. What diet should I follow?	52 · high 2 hidden, health kept
A5	Why was my payment declined? Card 4111 1111 1111 1111, expiry 12/28, OTP 482913.	99 · critical card + OTP removed
A5	Email me at priya.nair [at] gmail [dot] com and call +91-98450-12345.	email + phone hidden
B2	I'm Rahul Sharma. My PAN is PQRSX6789K and my UPI is rahul.sharma@okhdfcbank. Draft an email to HDFC support, IFSC HDFC0001234, asking why my refund is late.	55 · high 4 hidden
B3	Fix this code: const key = "sk-proj-Xy70m2Lp9Rt4Vb8Nc3Kd6Hf1Js5Gw0Ez"; fetch(url, { headers: { Authorization: key } })	71 · critical key removed
C6	My PAN is PQRSX6789K. Write one line confirming my PAN.	sent as «PAN_1»
C6	my password is Hunter@123	not sent

Downloads: github.com/Darshanj-dev/Mirage#download (extension, desktop) · github.com/NinadPandith/Mirage-CLI (CLI). Scores are from MIRAGE Core v1.0.0, checked 29 Sep 2026; the A5 disguised-email input was checked without the PAN, because a PAN split with spaces is not caught.