

MIRAGE · Presentation Script

What to say, how to say it, and what to show. About 10 minutes + Q&A. Replace [Name] with your team members.

TIME	SECTION	SPEAKER	SCORES FOR
0:00–1:00	1. The problem (hook)	Darshan	Social impact
1:00–2:00	2. What MIRAGE is	Darshan	Innovation
2:00–4:30	3. Demo: browser extension	Darshan	Implementation
4:30–6:30	4. Demo: desktop companion	Darshan	Implementation, innovation
6:30–8:30	5. Demo: CLI	Ninad	Implementation, teamwork
8:30–9:30	6. How it works + what's new	[Name]	Innovation
9:30–10:00	7. Impact, team, close	Darshan	Social impact, teamwork

How to deliver: start with the problem, not the tech. Talk to the judges, not the screen. After every demo step, say **one sentence about what just happened** ("ChatGPT never saw the PAN"). Slow down on the ★ moments. If something fails, say "let me show it another way" and move to the backup. Never pretend.

1 · The problem

0:00–1:00 · Darshan

SAY

"Show of hands: who has pasted something personal into ChatGPT? A bank letter, a resume, a screenshot of an error?"

Millions of Indians now use ChatGPT for bank letters, KYC emails and job applications, and they paste their **Aadhaar, PAN and phone number** straight in. Developers paste logs and code with **API keys and passwords**. In 2023, Samsung engineers leaked internal source code this way, and Samsung had to ban these tools.

Once you press Send, that data is on someone else's server. You can't take it back. And India's new **Digital Personal Data Protection Act, 2023** makes protecting personal data a legal duty. People need AI, but they need it **safely**."

Show: title slide. **How:** pause after the question; let them think. Serious tone.

2 · What MIRAGE is

1:00–2:00 · Darshan

"MIRAGE is a **privacy firewall for AI**. It sits between you and the AI and checks every prompt **before** it is sent, on your own device.

Personal details like Aadhaar and PAN become **placeholders**, so the AI can still answer. Secrets like API keys and passwords are **removed completely**. And when the answer comes back, MIRAGE **puts your real details back**, only on your screen.

Most tools just block you. MIRAGE lets you keep using AI and keeps your data out of it. We built it three ways, because people use AI in three places: the **browser**, the **ChatGPT and Claude desktop apps**, and the **terminal**."

Show: the flow slide: "Aadhaar · PAN · API key → MIRAGE (on your device) → «AADHAAR_1» «PAN_1» «API_KEY_REMOVED» → AI".

3 · Demo: browser extension

2:00–4:30 · Darshan

Do: chatgpt.com, new chat. Paste **A1** (clean prompt) → it sends normally.

"First, a normal question. Nothing sensitive, so MIRAGE stays out of the way. No false alarms."

Do ★: paste **A2** (Priya: PAN, Aadhaar, phone, email) → press Enter → the review opens. Click **Review**, then **Protect & send**.

"Now Priya asks for a bank letter with her PAN, Aadhaar, phone and email. I press Enter, and MIRAGE **stops it**. It found 5 items, the risk is **100 out of 100**, and it tells me exactly why: every point is explained. Nothing has been sent yet.

I click Protect & send. Look at what ChatGPT received: «**PAN_1**», «**AADHAAR_1**». ChatGPT never saw the real numbers. But it can still write the letter. And when I hover over a placeholder in the reply, the real value appears, **only on this laptop**."

Do: paste **A3** (AWS key + password) → Protect & send.

"Secrets are different. An API key isn't hidden, it's **removed**, never stored, never put back. One leaked cloud key can cost a company lakhs."

Optional (if time): A4 health: "The medical detail stays so the answer is useful, but the name and date of birth are hidden."

4 · Demo: desktop companion

4:30–6:30 · Darshan

"But many people don't use the website. They use the **ChatGPT desktop app**. A browser extension can't see inside that app, so we built a **native desktop companion**."

Show: the menu bar shield and the green "**MIRAGE active**" badge above ChatGPT's message box.

"This badge shows MIRAGE is protecting this box. It reads only the message box of ChatGPT and Claude, through the same accessibility system screen readers use. **No keylogger, no screen recording**."

Do ★: type **B2** (Rahul: PAN, UPI, IFSC) and stop typing → the popup opens by itself. Press **Enter** → nothing happens. Click **Protect & Send**.

"As soon as I stop typing, the review opens, in about **0.15 seconds**. Watch: I press Enter, and **nothing is sent**. MIRAGE is holding the message. I choose Protect & Send, and the chat shows only placeholders."

Optional: B3 (API key in code) → "Secret · won't be sent". Mention Windows: "We also have a Windows version in beta, installed with one click."

5 · Demo: CLI

6:30–8:30 · Ninad

"I built the developer side. Developers paste **logs and code** into AI all day, and logs are full of secrets and customer data."

Do: `cat error.log → mirage mask error.log (stops) → mirage mask error.log --strip-secrets → mirage unmask reply.txt.`

"Here's a real-looking server log: an API key, a customer email, an Aadhaar number, a phone number. First, MIRAGE **refuses** because there's a secret. With `strip-secrets` it removes the key and hides the rest: this is safe to paste into any AI. When the AI answers, **unmask** puts the real details back."

Do (if Wi-Fi is good): `mirage chat claude → "My PAN is PQRSX6789K..." → point at the sent › line, then the real PAN in the answer.`

"And we can **prove** it: this is Claude's own saved history on this laptop. The real PAN appears **zero** times. Claude only ever saw «PAN_1». And it's the **same detection engine** as the browser and desktop apps."

6 · How it works and what's new

8:30–9:30 · [Name]

"Under the hood, it's **one detection engine** written in TypeScript, and the same engine runs in all three products.

It doesn't just match patterns. It **verifies** them: Aadhaar with the **Verhoeff checksum**, card numbers with **Luhn**, PAN with its format rules, so random numbers don't cause false alarms. It measures **entropy** to spot random-looking API keys, and it reads **context**: words like 'DOB' or 'my name is'. It even catches tricks like 'gmail [dot] com' or hidden characters.

Then it gives a **risk score from 0 to 100** and explains every point. On our test set of 181 prompts, precision is 1.0 and recall 0.996, in 0.05 milliseconds per prompt. We wrote that test set ourselves, so an independent test is our next step.

What's new: **India-first detection** with real checksums, **reversible placeholders** so the AI still answers, and **interception inside native desktop apps**, not just websites. And there's **no server**: real values are encrypted on your device and deleted after 24 hours."

Show: the tech stack slide: one engine → Extension (WXT + React) · Desktop (Swift / .NET) · CLI (Node.js).

7 · Impact, team, close

9:30–10:00 · Darshan

"Who is this for? **Students** writing applications, **employees** using AI at work, **small businesses** handling customer data, and **developers**. Anyone who uses AI and has something to protect.

We're a team of [N]: I built the browser extension and the desktop companion, Ninad built the CLI, [Name] did [part]. We worked on **one shared engine**, so every improvement reaches all three products at once.

MIRAGE: **use AI without giving away who you are**. Thank you. We're happy to take questions."

Show: final slide with the GitHub links and a QR code for github.com/Darshanj-dev/Mirage. Smile, stop talking, wait.

Q&A · short, honest answers

Anyone

IF THEY ASK	SAY
Does data go to your server?	"No. There's no MIRAGE server. Detection runs on the device; only the protected prompt goes to the AI."

Why not use an AI model to detect?	"Speed and privacy: rules with checksums run in 0.05 ms, offline, and every result is explainable. A model would need a server or a big download."
What does it miss?	"Names only with context like 'my name is'; a PAN split with spaces; non-English text. That's why the review shows exactly what will be hidden, so the user stays in control."
Isn't the desktop app a keylogger?	"No. It only reads ChatGPT's and Claude's message box while they're open, only reacts to Enter and the Send button, never records the screen and never logs text."
How is it different from existing tools?	"Enterprise tools block and report to admins. MIRAGE is for individuals: it hides details so the AI still works, restores them in the reply, and understands Indian IDs with real checksums."
Business model / next steps?	"Free for individuals. Next: a blind accuracy test, the Chrome Web Store, testing the Windows app with ChatGPT, and a team version for companies under the DPDP Act."
Is Windows ready?	"It's a beta: built and tested automatically on Windows, but we haven't tried it with the ChatGPT app on Windows yet."

Before you go on: rehearse the whole script twice with a timer. Every speaker knows the ★ moments. Inputs to copy are in **MIRAGE-Demo-Inputs.pdf**; the step-by-step demo with backups is in **MIRAGE-Demo-Script.pdf**.